

CS3811 - High Performance Computing and Big Data Lab

Lab 1

Name: M K Lokesh Kumar

Registration No.: 2201113026

Class: Cyber Security(Semester 5)

Experiment 1

Objective

To estimate the value of π using a randomised algorithm technique(Monte Carlo), and plot 2 performance graphs:

- Graph 1 - Time taken by the code vs Number of samples
- Graph 2 - Error vs number of samples

Code

Written in C++.

```
#include <iostream>
#include <cstdlib>
#include <ctime>
#include <cmath>
#include <chrono>
#include <vector>

double estimatePi(int numSamples) {
    int pointsInsideCircle = 0;

    for (int i = 0; i < numSamples; i++) {
        double x = static_cast<double>(rand()) / RAND_MAX;
        double y = static_cast<double>(rand()) / RAND_MAX;

        if (x * x + y * y <= 1) {
            pointsInsideCircle++;
        }
    }

    return 4.0 * pointsInsideCircle / numSamples;
}
```

```

int main() {
    srand(static_cast<unsigned int>(time(0)));
    const int maxSamples = 1000000;
    const int step = 10000;

    std::vector<int> sampleSizes;
    std::vector<double> times;
    std::vector<double> errors;
    std::vector<double> pi;

    for (int numSamples = step; numSamples <= maxSamples; numSamples +=
step) {
        auto start = std::chrono::high_resolution_clock::now();
        double piEstimate = estimatePi(numSamples);
        auto end = std::chrono::high_resolution_clock::now();

        std::chrono::duration<double> duration = end - start;
        double timeTaken = duration.count();
        double error = std::abs(M_PI - piEstimate);

        sampleSizes.push_back(numSamples);
        times.push_back(timeTaken);
        errors.push_back(error);
        pi.push_back(piEstimate);
    }

    std::cout << "Samples\tTime\tError\tPi Estimate" << std::endl;
    for (size_t i = 0; i < sampleSizes.size(); ++i) {
        std::cout << sampleSizes[i] << "\t" << times[i] << "\t" <<
errors[i] << "\t" << pi[i] << std::endl;
    }

    return 0;
}

```

Gnuplot script for plotting performance graphs 1 and 2

```

set terminal pngcairo enhanced font 'Verdana,12'

set output 'ex_1_time_vs_samples.png'
set title 'Time Taken vs Number of Samples'
set xlabel 'Number of Samples'
set ylabel 'Time Taken (seconds)'
plot 'data1.dat' using 1:2 with linespoints title 'Time Taken'

set output 'ex_1_error_vs_samples.png'
set title 'Error in Estimation of  $\pi$  vs Number of Samples'
set xlabel 'Number of Samples'
set ylabel 'Error'

```

```
plot 'data1.dat' using 1:3 with linespoints title 'Error'
```

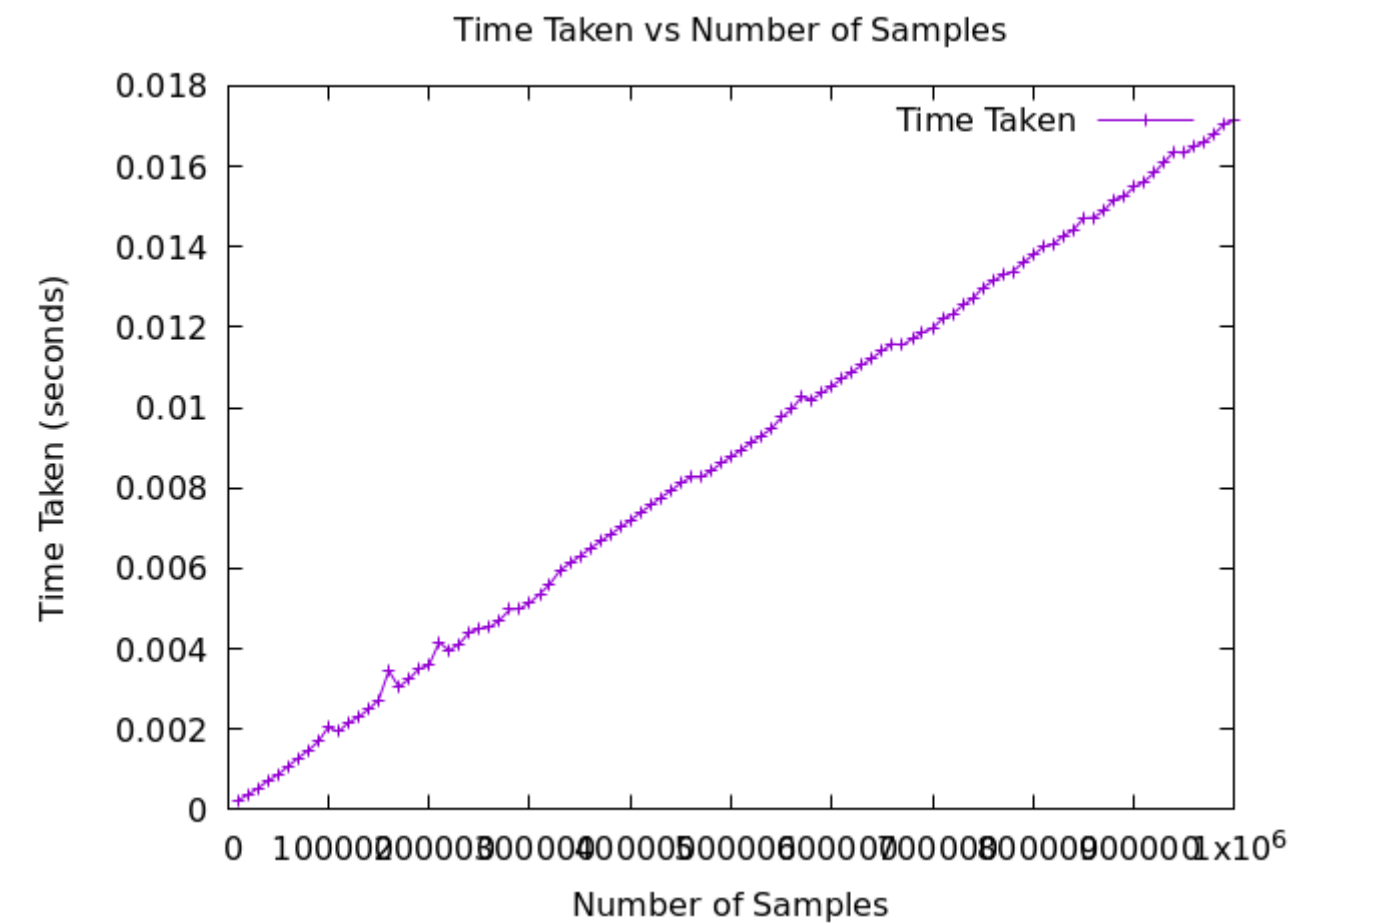
Output

Samples	Time	Error	Pi Estimate
10000	0.000278414	0.000807346	3.1424
20000	0.000610885	0.0120073	3.1536
30000	0.000860183	0.00239265	3.1392
40000	0.00113147	0.00409265	3.1375
50000	0.00138829	0.00952735	3.15112
60000	0.00167266	0.00754068	3.14913
70000	0.001954	0.00416408	3.13743
80000	0.00227491	0.00649265	3.1351
90000	0.00258685	0.00227401	3.14387
100000	0.00278274	0.00795265	3.13364
110000	0.00321252	0.00355629	3.13804
120000	0.003246	0.000592654	3.141
130000	0.00361894	0.00728496	3.13431
140000	0.00403375	0.0053502	3.14694
150000	0.0042086	0.00200735	3.1436
160000	0.0045163	0.000967654	3.14062
170000	0.00475619	0.00530146	3.14689
180000	0.00502096	0.0019629	3.14356
190000	0.00533114	0.0064284	3.14802
200000	0.00536131	0.0100127	3.13158
210000	0.00588245	0.00242639	3.14402
220000	0.00611366	0.000992654	3.1406
230000	0.00612452	0.001697	3.1399
240000	0.00639209	0.00817401	3.14977
250000	0.0069648	0.000952654	3.14064
260000	0.00720612	0.00582342	3.13577
270000	0.00737869	7.4684e-06	3.14159
280000	0.00773827	0.00129306	3.14289
290000	0.00797314	0.000986657	3.14258
300000	0.00817643	0.00497932	3.13661
310000	0.00839234	0.000147492	3.14145
320000	0.00892622	0.000344846	3.14194
330000	0.0090751	0.00348613	3.14508
340000	0.00930424	0.000651477	3.14094
350000	0.0091494	0.0031502	3.14474
360000	0.0104005	0.000851791	3.14244
370000	0.0108487	0.00149924	3.14309
380000	0.0107987	0.00294002	3.13865
390000	0.0111005	0.000109911	3.1417
400000	0.0113232	0.00199265	3.1396
410000	0.0115427	0.000529239	3.14106
420000	0.0117598	0.00260735	3.1442
430000	0.0123287	0.00251897	3.14411
440000	0.0124723	0.00229826	3.14389
450000	0.0127661	0.00323401	3.14483
460000	0.0129028	0.00307961	3.13851
470000	0.0132835	0.0012267	3.14037
480000	0.0136069	0.00117599	3.14042
490000	0.0136849	0.00134776	3.14024
500000	0.013914	0.00436735	3.14596
510000	0.0142011	0.000736758	3.14233
520000	0.0145825	0.000338116	3.14193
530000	0.0145695	0.00238511	3.13921
540000	0.0151355	0.00126661	3.14286
550000	0.0151262	0.00174538	3.13985


560000	0.0147679	7.83679e-05	3.14151
570000	0.0152571	0.00068739	3.14091
580000	0.0153061	0.000992654	3.1406
590000	0.0151568	0.00258362	3.14418
600000	0.0159686	0.00288735	3.14448
610000	0.016713	0.000792654	3.1408
620000	0.0169281	0.00396219	3.14555
630000	0.01734	0.000172426	3.14177
640000	0.0171854	3.23464e-05	3.14162
650000	0.0170866	0.000552654	3.14104
660000	0.017753	0.00178659	3.13981
670000	0.0176369	0.00332078	3.14491
680000	0.0187227	0.000260288	3.14185
690000	0.0196701	0.00139285	3.14299
700000	0.0187286	0.000478368	3.14111
710000	0.0186236	0.00233974	3.14393
720000	0.0192235	0.00249624	3.14409
730000	0.0188429	0.00303375	3.13856
740000	0.0195392	0.00281428	3.13878
750000	0.0194673	0.000167346	3.14176
760000	0.0196735	0.00276524	3.14436
770000	0.0208195	0.000693061	3.14229
780000	0.0208039	0.000920167	3.14251
790000	0.0215895	0.000534426	3.14106
800000	0.0221681	0.000697346	3.14229
810000	0.0222894	0.000891419	3.1407
820000	0.022644	0.00327558	3.13832
830000	0.0226114	0.00129386	3.1403
840000	0.0228508	0.000745035	3.14085
850000	0.0235773	0.000717359	3.14088
860000	0.0240288	0.000998044	3.14259
870000	0.0243733	0.000406447	3.14119
880000	0.0245022	0.00233007	3.14392
890000	0.0246128	0.000325238	3.14127
900000	0.0250051	0.00139265	3.1402
910000	0.0258083	0.00307397	3.13852
920000	0.025845	0.00165352	3.13994
930000	0.0259343	0.000861471	3.14073
940000	0.0261734	0.000190325	3.14178
950000	0.0268085	0.00119686	3.1404
960000	0.0274866	0.00050932	3.14108
970000	0.0267407	0.000363788	3.14123
980000	0.0265272	4.57148e-05	3.14155
990000	0.0265179	0.000512397	3.14211
1000000	0.0281924	0.000848654	3.14074

Performance Graphs

- Time taken by the graph vs Number of samples



- Error vs Number of samples



Experiment 2

Objective

Write a C/C++ program to perform matrix-vector multiplication using 2 different methods.

$$b := Ax$$

In the above, $b = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_N \end{bmatrix}$, $A = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1N} \\ a_{21} & a_{22} & \cdots & a_{2N} \\ \vdots & \vdots & \ddots & \vdots \\ a_{N1} & a_{N2} & \cdots & a_{NN} \end{bmatrix}$ and $x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_N \end{bmatrix}$

Method 1

$$b_i := \sum_{j=1}^N a_{ij}x_j, \forall i \in \{1, 2, \dots, N\}$$

Method 2

$$\begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_N \end{bmatrix} = x_1 \times \begin{bmatrix} a_{11} \\ a_{21} \\ \vdots \\ a_{N1} \end{bmatrix} + x_2 \times \begin{bmatrix} a_{12} \\ a_{22} \\ \vdots \\ a_{N2} \end{bmatrix} + \cdots + x_N \times \begin{bmatrix} a_{1N} \\ a_{2N} \\ \vdots \\ a_{NN} \end{bmatrix}$$

Also, plot 2 performance graphs:

- Graph 1 - Time taken by Method 1 and Method 2 vs N
- Graph 2 - FLOPs vs N

Code

Written in C++.

```
#include <iostream>
#include <vector>
#include <chrono>

void method1(const std::vector<std::vector<double>>& A, const
std::vector<double>& x, std::vector<double>& b) {
    int N = A.size();
    for (int i = 0; i < N; i++) {
        b[i] = 0;
        for (int j = 0; j < N; j++) {
            b[i] += A[i][j] * x[j];
        }
    }
}
```

```

}

void method2(const std::vector<std::vector<double>>& A, const
std::vector<double>& x, std::vector<double>& b) {
    int N = A.size();
    for (int i = 0; i < N; i++) {
        b[i] = 0;
        for (int j = 0; j < N; j++) {
            b[i] += A[j][i] * x[i];
        }
    }
}

int main() {
    const std::vector<int> sizes = {128, 512, 1024, 2048, 4096, 8192};

    std::cout << "N" << "\t" << "Time taken by method 1" << "\t" << "Time
taken by method 2" << "\t" << "FLOPs" << std::endl;

    for (int N : sizes) {
        std::vector<std::vector<double>> A(N, std::vector<double>(N));
        std::vector<double> x(N);
        std::vector<double> b1(N, 0);
        std::vector<double> b2(N, 0);

        for (int i = 0; i < N; i++) {
            for (int j = 0; j < N; j++) {
                A[i][j] = static_cast<double>(rand()) / RAND_MAX;
            }
            x[i] = static_cast<double>(rand()) / RAND_MAX;
        }

        auto start1 = std::chrono::high_resolution_clock::now();
        method1(A, x, b1);
        auto end1 = std::chrono::high_resolution_clock::now();
        std::chrono::duration<double> duration1 = end1 - start1;
        double timeMethod1 = duration1.count();

        auto start2 = std::chrono::high_resolution_clock::now();
        method2(A, x, b2);
        auto end2 = std::chrono::high_resolution_clock::now();
        std::chrono::duration<double> duration2 = end2 - start2;
        double timeMethod2 = duration2.count();

        double flops = 2.0 * N * N;

        std::cout << N << "\t" << timeMethod1 << "\t" << timeMethod2 <<
"\t" << flops << std::endl;
    }

    return 0;
}

```


Gnuplot script for plotting performance graphs 1 and 2

```
set terminal pngcairo enhanced font 'Verdana,12'

set output 'ex_2_time_vs_N.png'
set title 'Time Taken by Method 1 and Method 2 vs N'
set xlabel 'N (Matrix Size)'
set ylabel 'Time taken (seconds)'
plot 'data2.dat' using 1:2 with linespoints title 'Method 1', \
      'data2.dat' using 1:3 with linespoints title 'Method 2'

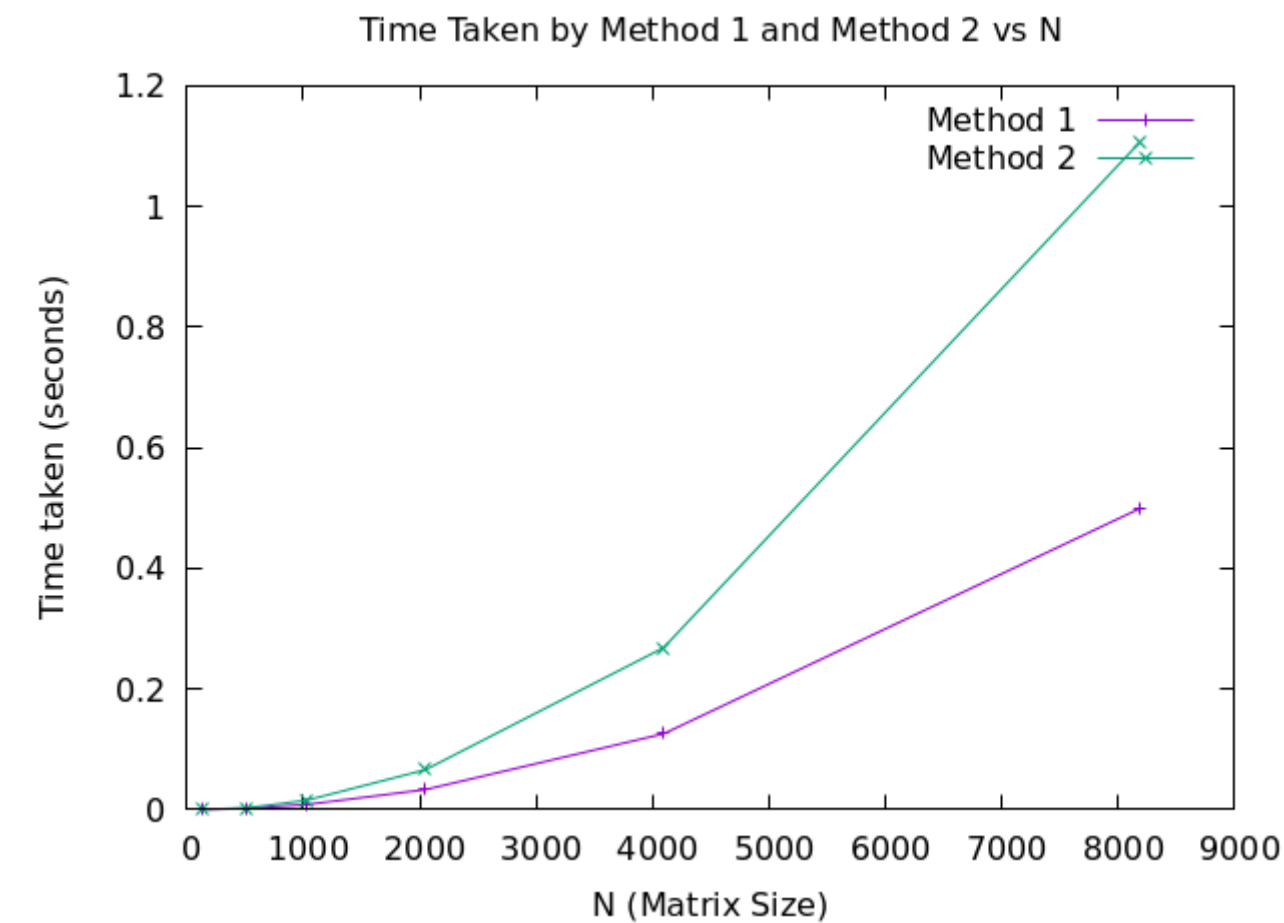
set output 'ex_2_flops_vs_N.png'
set title 'FLOPs vs N'
set xlabel 'N (Matrix Size)'
set ylabel 'FLOPs'
plot 'data2.dat' using 1:4 with linespoints title 'FLOPs'
```

Output

N	Time taken by method 1	Time taken by method 2	FLOPs
128	8.4685e-05	8.7317e-05	32768
512	0.0015207	0.00206153	524288
1024	0.00668042	0.0124618	2.09715e+06
2048	0.0248162	0.0464221	8.38861e+06
4096	0.0923269	0.178096	3.35544e+07
8192	0.402208	0.961143	1.34218e+08

Performance Graphs

- Graph 1 - Time taken by Method 1 and Method 2 vs N



- Graph 2 - FLOPs vs N

