

CS3811 - High Performance Computing and Big Data Lab

Lab 6

Name: M K Lokesh Kumar

Registration No.: 2201113026

Class: Cyber Security(Semester 5)

Experiment 1

Objective

To perform sparse matrix computations and plot a performance graph comparing the execution time against the number of processing elements used for parallelization.

Code

Written in C++.

The matrix used here was downloaded from the Matrix Market, and has 1138 rows, 1138 columns, and 2596 non-zero values in it. It can be found here: [Link to Matrix](#)

```
#include <iostream>
#include <vector>
#include <fstream>
#include <omp.h>
#include <chrono>

using namespace std;

struct CSRMatrix {
    int rows, cols;
    vector<int> row_ptr, col_idx;
    vector<double> vals;
};

CSRMatrix read_mtx(const string &filename) {
    ifstream infile(filename);

    CSRMatrix matrix;
    int nz;
    infile >> matrix.rows >> matrix.cols >> nz;

    matrix.row_ptr.resize(matrix.rows + 1);
```

```
matrix.col_idx.resize(nz);
matrix.vals.resize(nz);

int row, col;
double val;

vector<int> row_count(matrix.rows + 1, 0);

for (int i = 0; i < nz; i++) {
    infile >> row >> col >> val;
    row--;
    col--;
    matrix.col_idx[i] = col;
    matrix.vals[i] = val;
    row_count[row + 1]++;
}

for (int i = 1; i <= matrix.rows; i++) {
    matrix.row_ptr[i] = matrix.row_ptr[i - 1] + row_count[i];
}

return matrix;
}

void spmv(const CSRMatrix &matrix, const vector<double> &v, vector<double>
&res, int threads) {
    #pragma omp parallel for num_threads(threads)
    for (int i = 0; i < matrix.rows; i++) {
        double sum = 0.0;
        for (int j = matrix.row_ptr[i]; j < matrix.row_ptr[i + 1]; j++) {
            sum += matrix.vals[j] * v[matrix.col_idx[j]];
        }
        res[i] = sum;
    }
}

int main() {
    ofstream outfile("data_1_0.dat");
    outfile << "Num of threads" << "\t" << "Time in seconds" << endl;

    CSRMatrix matrix = read_mtx("1138_bus.mtx");

    vector<double> v(matrix.cols, 1.0);
    vector<double> res(matrix.rows, 0.0);

    vector<int> numThreads = {1, 2, 4, 8};

    for (int i = 0; i < numThreads.size(); i++) {
        double start = omp_get_wtime();
        spmv(matrix, v, res, numThreads[i]);
        double duration = omp_get_wtime() - start;
        outfile << numThreads[i] << "\t" << duration << endl;
    }
}
```

```
    return 0;  
}
```

Gnuplot script for plotting performance graph

```
set terminal pngcairo enhanced font 'Arial,10'  
set output 'spmv_plot_1_0.png'  
  
set title "Execution Time vs. Number of Threads"  
set xlabel "Number of Threads"  
set ylabel "Time (seconds)"  
set grid  
  
set logscale y 10  
set xrange [0:9]  
  
plot "data_1_0.dat" using 1:2 with linespoints \  
    title "Execution Time" lw 2 pt 7 lc rgb "blue"
```

Output

```
Num of threads  Time in seconds  
1      3.4269e-05  
2      0.000154159  
4      5.8178e-05  
8      7.4634e-05
```

Performance Graphs

- Execution time vs Number of processing elements

